

## TMA4280 - Superdatamaskiner

### Exercise 4

Eirik Mildestveit Hammerstad

I just want to say sorry if anything is not up to level of your expectation. I've caught the flu and 2 hours before delivery I managed to rm \* -rf on the wrong directory, so sorry if some of the code/report is not up to standards.

1)

The serial version can be found in serial.c. It outputs the following:

|              |                    |                           |
|--------------|--------------------|---------------------------|
| S - S(16)    | 1.5843465334449871 | Error: 0.0605875334032393 |
| S - S(32)    | 1.6141672628279242 | Error: 0.0307668040203022 |
| S - S(64)    | 1.6294305014088870 | Error: 0.0155035654393394 |
| S - S(128)   | 1.6371520049544612 | Error: 0.0077820618937652 |
| S - S(256)   | 1.6410354363086801 | Error: 0.0038986305395463 |
| S - S(512)   | 1.6429828479550965 | Error: 0.0019512188931299 |
| S - S(1024)  | 1.6439579810301654 | Error: 0.0009760858180610 |
| S - S(2048)  | 1.6444459047881168 | Error: 0.0004881620601096 |
| S - S(4096)  | 1.6446899560231332 | Error: 0.0002441108250932 |
| S - S(8192)  | 1.6448120039860050 | Error: 0.0001220628622214 |
| S - S(16384) | 1.6448730335545856 | Error: 0.0000610332936408 |

The serial code is quite straight forward and naive: For  $n = 2^k$ ,  $k = 4, \dots, 16$  i do the following: I create a vector where all elements are  $1^{-2}$ , and sum the elements in the vector. The vector\_t struct can be found in common.h.

2)

The parallel version with shared memory can be found in parallel-shared.c. It outputs the following:

|              |                    |                           |
|--------------|--------------------|---------------------------|
| S - S(16)    | 1.5843465334449871 | Error: 0.0605875334032393 |
| S - S(32)    | 1.6141672628279242 | Error: 0.0307668040203022 |
| S - S(64)    | 1.6294305014088870 | Error: 0.0155035654393394 |
| S - S(128)   | 1.6371520049544612 | Error: 0.0077820618937652 |
| S - S(256)   | 1.6410354363086803 | Error: 0.0038986305395461 |
| S - S(512)   | 1.6429828479550954 | Error: 0.0019512188931310 |
| S - S(1024)  | 1.6439579810301632 | Error: 0.0009760858180632 |
| S - S(2048)  | 1.6444459047881148 | Error: 0.0004881620601116 |
| S - S(4096)  | 1.6446899560231238 | Error: 0.0002441108251026 |
| S - S(8192)  | 1.6448120039860075 | Error: 0.0001220628622189 |
| S - S(16384) | 1.6448730335545798 | Error: 0.0000610332936466 |

The parallel version with shared memory works much like the serial, with the exception that the summation is done in parallel (with a reduction on the sum variable).

3)

The parallel version with distributed memory can be found in parallel-distributed.c. It outputs the following:

|              |                    |                           |
|--------------|--------------------|---------------------------|
| S - S(16)    | 1.5843465334449871 | Error: 0.0605875334032393 |
| S - S(32)    | 1.6141672628279242 | Error: 0.0307668040203022 |
| S - S(64)    | 1.6294305014088870 | Error: 0.0155035654393394 |
| S - S(128)   | 1.6371520049544612 | Error: 0.0077820618937652 |
| S - S(256)   | 1.6410354363086801 | Error: 0.0038986305395463 |
| S - S(512)   | 1.6429828479550965 | Error: 0.0019512188931299 |
| S - S(1024)  | 1.6439579810301654 | Error: 0.0009760858180610 |
| S - S(2048)  | 1.6444459047881168 | Error: 0.0004881620601096 |
| S - S(4096)  | 1.6446899560231332 | Error: 0.0002441108250932 |
| S - S(8192)  | 1.6448120039860050 | Error: 0.0001220628622214 |
| S - S(16384) | 1.6448730335545856 | Error: 0.0000610332936408 |

The parallel version with distributed memory had to be rewritten after an accidental removal of some files. The first processor creates a vector and splits it up to the other processors, which sum the elements and return the answer, which is added together on the first processor.

4)

Necessary/convenient MPI calls:

- MPI\_Init
- MPI\_Comm\_size
- MPI\_Comm\_rank
- MPI\_Get\_processor\_name
- MPI\_Send
- MPI\_Recv
- MPI\_Finalize

5)

If we compare the results from the different programs, they are identical for  $n=4$ , but for  $n=16$ , then parallel yield a better result. I believe this has to do with the representation of floating point numbers. Once we start losing precision when adding new elements, the sum gets influenced, something which happens at a stronger rate in the sequential version.

6)

In the sequential version, we generate and hold everything in one shared memory, while the parallel distributed version creates everything in a single process, and copies parts of the vector to other processes, which would use more ram.

7)

I assume that it would create two floating point operations per item in `vector_t->data`, one to sum them, and one for the minus, so  $4n$  where  $n$  is elements.

Only one processor creates the vector, which is a significant part of the exercise, so it is not loadbalanced.

8)

No, I believe we would get better results with a dynamic programming implementation.